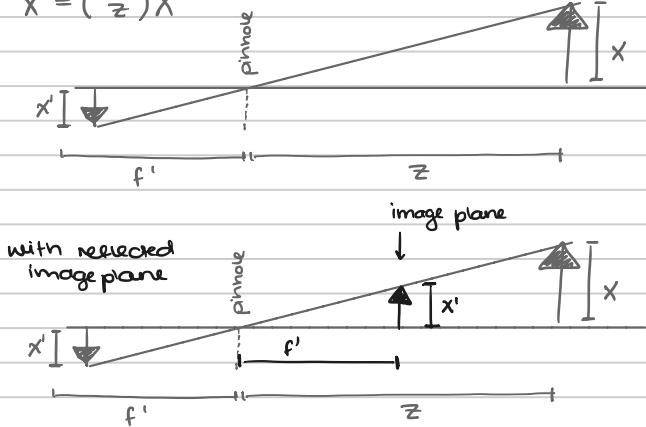


LEC1 - Image formation & basics

Lec 1-1

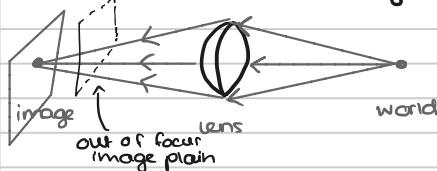
Camera obscura : pinhole mirrors upside-down but not left-to-right like mirror!

$$x' = \left(\frac{f'}{z}\right)x$$



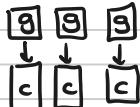
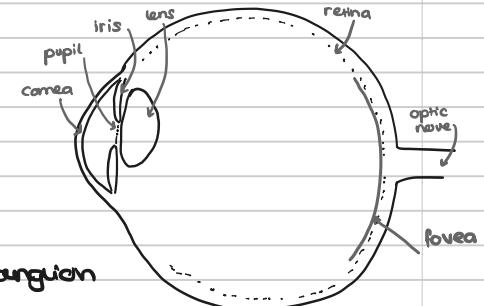
- distant objects are smaller
- Flight point = vanishing point lies on the horizon parallel lines meet
- each set of parallel lines meet at different vanishing points
- perspective projections:
 - point → point
 - line → line / point
 - planes → images / line
 - angles mostly not preserved

- pinhole camera limitations: too large → blurring, too small → diffraction → blur
 - no focus possible.
 - + too dark images
- invention of lenses: get more light & focus it

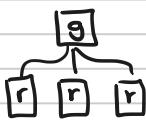


- lens: blurring if image plain out of focus
- today: lens systems ~ 15 elements
- human eye: our vision affects a lot how we tackle ip problems:

- iris like camera; focusing by changing lens shape
- Retina: cones (Zapfen): high resolution, precision
rod (Stäbchen): for low light, moves, much smaller than cones
- Fovea: high resolution, mostly cones



- 1 ganglion per cone, several rods per ganglion
- C: every photon separately registered
- r: highly sensitive to motion



- 3 coln types respond to blue, green, red

- CCD cameras (→ Bayer filter) inspired by cones
we perceive more green

$$\bullet I_c(x, y); c = \{0, \dots, 2\} \quad I_c = \{I_0, I_1, I_2\} \stackrel{\wedge}{=} \text{red, green, blue}$$

$$I(x, y) = \begin{matrix} I(0, 0) & \dots & I(0, M) \\ \dots & \dots & \dots \\ I(N, 0) & \dots & I(N, M) \end{matrix}$$

illumination

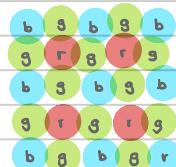
- Luminosity: amount of light, brightness
- $\ell: X \times Y \rightarrow \mathbb{R}^+$

$$\text{Reflectance } r: X \times Y \rightarrow (0, 1) \quad 0 \leq r(x, y) \leq 1$$

$\hookrightarrow I(x, y) \in \{0, \dots, +\infty\} \mathbb{R}^+$: enormous range! (and strong simplification)

\hookrightarrow Quantization: to represent values of I e.g. 2 values per pixel $\in (0, 1) \rightarrow 2^2$ possibilities 00 01 10 11

- Sampling: to map real-value signals to discrete space
some dist between (0, 1) and (1, 2) ?



Bayer filter
25% red/blue
50% green

- Sampling:  vs.  high vs. low resolution, how many samples?

- Quantization:  steps don't have to have same size! Stems makes sense to give all highly dark pixel the same value "how many"

• Image values $\in \{0, 255\}$ $0 = \text{black}, 255 = \text{white} \rightarrow 256 \text{ values}$
 color image: 3 channels $I = \{I_1, I_2, I_3\}$

- Tessellation: picture elements = pixels, how to separate space

 vs  vs semi-regular tessellation
 \rightarrow interior summe := 360°
 covers entire space!

- Image model $I = l \cdot r$

- 1. Integrate brightness over specified region.
 whole world: $p(x, y)$ is "everywhere"

$$I(u, v) = \iint_{x, y} l(x, y) p(x-u, y-v) dx dy$$



exact value of $I(u, v)$?

$$p(x, y) = \begin{cases} 1 & \text{if } x, y \text{ are in pixel } u, v \\ 0 & \text{else} \end{cases}$$

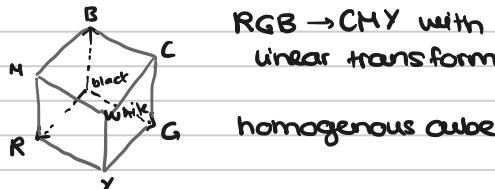
\uparrow Dirac function (δ) 'indicator function'

- visible range: 400 - 700 nm

additive (projector) vs. subtractive (printer) color mixing

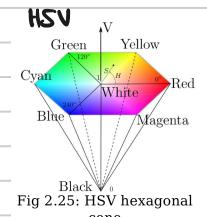
\uparrow RGB, black at beginning
 \uparrow at beginning it's white, add all colors \rightarrow black CMY

- RGB color model vs CMY color model $\in \{0, 255\}$



RGB \rightarrow CMY with rotation
 linear transformation

$$\begin{bmatrix} C \\ M \\ Y \end{bmatrix} = \begin{bmatrix} 1-R \\ 1-G \\ 1-B \end{bmatrix}$$



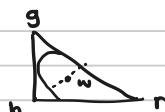
- HSV: Hue, Saturation, Value: hexagon shaped, edges are pure colors, move inwards = mix colors

more close to human perception

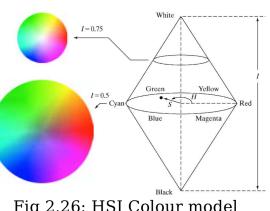
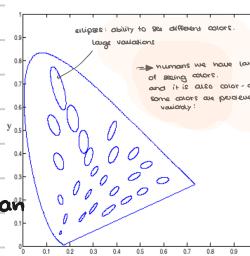
= HSB ; no linear transformation to RGB $\sim H$ captures variations

- HSI: Hue, Saturation, Intensity = HSL ; color hexone / sphere
 nonlinear transformation to RGB

- CIE: white in middle.
 RGB is small part of it.



- McAdam ellipses: ability to differentiate colors is limited. / large variation among human and color-dependent



- Adaptation phenomena: color perception depends on spatial contrast & previously seen stuff (e.g. walk in on bright day)

\hookrightarrow be cautious with colors in algos! Stop sign is red \rightarrow but how to see it in shadow?

Distances

- can't be negative
- identity $d(x, x) = 0$
- symmetry $d(x, y) = d(y, x)$
- triangle inequality $d(x, z) \leq d(x, y) + d(y, z)$

x — y — z vs x — y — z

} most important for sip

- L_1 -norm: "city block" $\sum_{i=1}^n |x_i - y_i|$



L^2 -norm: euclidean norm $\sqrt{\sum_{i=1}^n (x_i - y_i)^2}$

L^p -norm: $(p \rightarrow \infty) \lim_{p \rightarrow \infty} \left(\sum_{i=1}^n |x_i - y_i|^p \right)^{\frac{1}{p}} = \max_{i=1 \dots n} \{|x_1 - y_1|, |x_2 - y_2|, \dots, |x_n - y_n|\}$
good where eucl. dist explodes

- Chamfer dist: (e.g. with L_1 -norm)
- build image such that $D(x, y) \in \{0, +\infty\}$
if value 0, it's 0 and else $+\infty$ (or a high number) for all $\forall x, y$

$$\begin{array}{|c|c|c|} \hline 0 & 1 & 0 & 0 \\ \hline 1 & 0 & 1 & 2 \\ \hline 0 & 1 & 2 & 3 \\ \hline \end{array} \quad I$$

$$\begin{array}{|c|c|c|} \hline 0 & 1 & 0 & 0 \\ \hline 1 & 0 & 1 & 2 \\ \hline 0 & 1 & 2 & 3 \\ \hline \end{array} \quad D$$

AL	AL	BR
AL	SK	BR
AL	BR	BR

- for $(x = 0 \dots N-1)$ left \rightarrow right
- for $(y = 0 \dots N-1)$ top \rightarrow bottom

$$D(x, y) = \min_{x, y \in AL} \{ D(x, y), D(x, y') + L^1 d((x, y), (x, y')) \}$$

e.g. 0 dist. zu (x', y') + dessen Wert
war schon bereit

- do the same for $x \in BR$

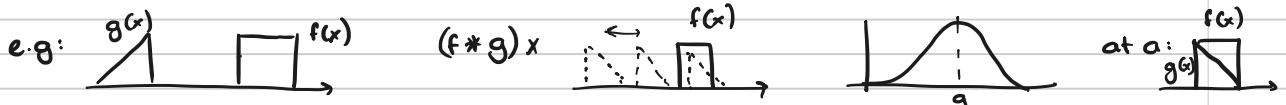
In words \sim go to every pixel, check above left of it, $O(\text{chamfer}) \sim N \cdot N \cdot 8$

LEC 2 - Image formation & basics (2)

L2 - 1

- $f, g : 1D$ continuous functions

- convolution : $(f * g)(x) = \int_{-\infty}^{\infty} f(\tau) g(x-\tau) d\tau = \int f(x-\tau) g(\tau) d\tau$



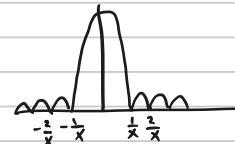
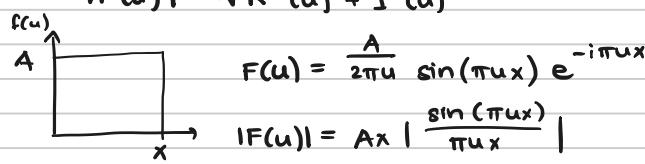
- Fourier decomposition : $F(u) = \int f(x) e^{-2\pi i ux} dx = \frac{1}{N} \sum_{n=0}^{N-1} f(x) e^{-\frac{i 2\pi u x}{N}}$

$$f(x) = \int F(u) e^{2\pi i ux} du ; \quad f \text{ is continuous}$$

if it's integrable

as $F(u) = \text{real}(u) + i \text{img}(u) = |F(u)| \cdot e^{i\phi(u)}$ $\phi(u) = \tan^{-1} I(u)/R(u)$

$$|F(u)| = \sqrt{R^2(u) + I^2(u)}$$



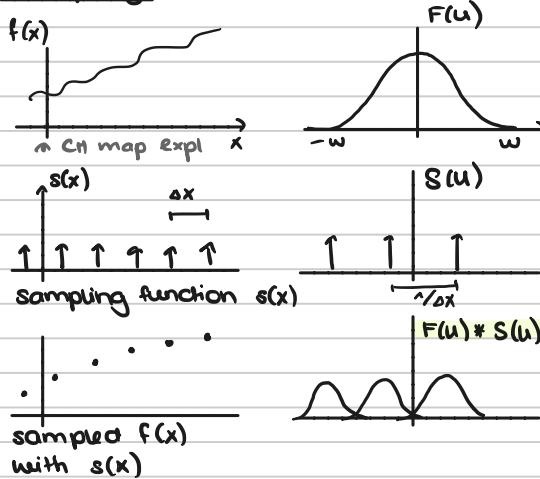
FFT algo in most processors, R, matlab, ...

- properties : $F(f+g) = F(f) + F(g)$ uniquely separable
 $F(f \cdot g) \neq F(f) \cdot F(g)$

• a, b : $a \cdot f(x) = a F(u)$
 $f(a \cdot x) = \frac{1}{a} F(u/a)$

$f(x) * g(x) = F(u) \cdot F(v)$	important!
$f(x) \cdot g(x) = F(u) * F(v)$	

Sampling



Aliasing : if $F(u) * S(u)$ overlap :

"duplication of signal"
 c : center of overlap
 $\text{every } \frac{1}{2\Delta x} = u$

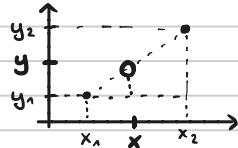
reduce overlap by getting more samples!
narrow sampling function $s(x) \rightarrow$ wide $S(u)$
→ we can't shrink an image by taking every second pixel → aliasing!
characteristic errors: small phenomena look bigger, fast look slower
e.g. wheels roll wrong way in TV

Nyquist theorem : $\Delta x \leq \frac{1}{2w}$
to recover the signal

Interpolation: a form of sampling
up-/downsampling possible

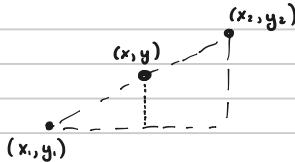
$$5 \boxed{} \rightarrow 3 \boxed{} \quad \text{or} \quad 10 \boxed{} \rightarrow \boxed{} 5$$

1D case:



$$y = ax + b$$

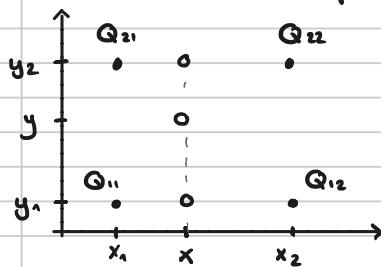
$$\frac{y - y_1}{x - x_1} = \frac{y_2 - y_1}{x_2 - x_1}$$



from triangle inequality

1D case is a linear operation, 2D not

2D case: Bilinear interpolation



$$\textcircled{1} \text{ at height } y_1: f(x, y_1) = \underbrace{\frac{x - x_1}{x_2 - x_1} f(Q_{11})}_{x_1} + \underbrace{\frac{x_2 - x}{x_2 - x_1} f(Q_{12})}_{x_2}$$

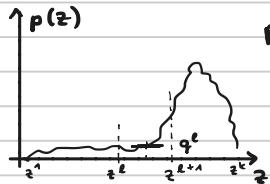
$$\textcircled{2} \text{ at height } y_2: f(x, y_2) = " f(Q_{21}) + " f(Q_{22})$$

\textcircled{3} interpolate in y-direction

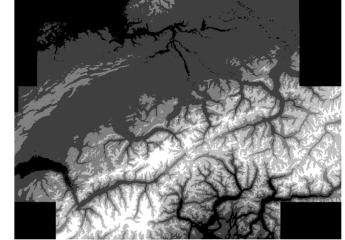
$$f(x, y) \approx \frac{y_2 - y}{y_2 - y_1} f(x, y_1) + \frac{y - y_1}{y_2 - y_1} f(x, y_2)$$

$$\text{for } x_2 - x_1 = 1 = y_2 - y_1: f(x, y) \approx [1-x \ x-1] \begin{bmatrix} Q_{11} & Q_{12} \\ Q_{21} & Q_{22} \end{bmatrix} \begin{bmatrix} 1-y \\ y-1 \end{bmatrix}$$

Quantization



prob. distribution of obs. value z . (left)
here: lots of high values \rightarrow distribute bins accordingly! not fixed bins



with $2^2 = 4$ grey values

Lloyd-Max Quantization: fixed amount of bins K
what interval to take from z^ℓ to $z^{\ell+1}$, q^ℓ : average height

- reduce the average quantization error with error function (loss function):

$$E = \sum_{\ell=1}^K \int_{z^\ell}^{z^{\ell+1}} (z - q^\ell)^2 p(z) dz$$

$$\bullet \text{ gradient descent: } \frac{\partial E}{\partial q^\ell} = -2 \int_{z^\ell}^{z^{\ell+1}} (z - q^\ell)^2 p(z) dz \stackrel{!}{=} 0$$

$$q^\ell = \frac{\int_{z^\ell}^{z^{\ell+1}} z p(z) dz}{\int_{z^\ell}^{z^{\ell+1}} p(z) dz}$$

$$\frac{\partial E}{\partial z^\ell} \stackrel{!}{=} 0 \Leftrightarrow z^\ell = \frac{z^\ell - z^{\ell+1}}{2}$$

Idea: first optimize q^ℓ with random z^ℓ , then optimise z^ℓ
loop until no change

Median cut algo for color images (3D, not possible for Lloyd-Max)

LEC 3 : Linear filtering

L3 - 1

$$\text{Correlation: } I'(x, y) = \sum_{i=-k}^k \sum_{j=-k}^k F(i, j) I(x+i, y+j)$$

$$\begin{aligned} \text{Convolution: } I'(x, y) &= "F(-i, -j)" \\ &= "I(x-i, y-j)" \end{aligned} \quad \text{"flipping"}$$

if $F(X, Y) = F(-X, Y) \rightarrow \text{Corr} = \text{Conv}$ (symmetric filter)

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} \rightarrow \text{no change} \quad \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix} \rightarrow \text{slide all pixels to the left} \quad \begin{bmatrix} 1/8 & 1/8 & 1/8 \\ 1/8 & 1/8 & 1/8 \\ 1/8 & 1/8 & 1/8 \end{bmatrix} \rightarrow \text{average filtering} = \text{blurring}$$

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \frac{1}{8} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \rightarrow \text{sharpening filter "laplacian"} \quad \begin{array}{c} 1 \\ -1 \\ -1 \end{array}$$

box filter: all weights equal, can show "lines" = artefacts

gaussian smoothing: no artefacts, looks like a fuzzy blob (what happens in reality)

$$1D: \mathcal{N}(\mu, \sigma) = \frac{1}{\pi} \exp\left(-\frac{1}{2\sigma^2}(x-\mu)^2\right)$$

$$2D: \mathcal{N}(x, \mu, \Sigma) = \frac{1}{2\pi} \exp\left(-\frac{1}{2} (x-\mu)^T \Sigma^{-1} (x-\mu)\right)$$

$$\Sigma: \text{covariance matrix} \in \mathbb{R}^{2 \times 2} \text{ if } \Sigma = \begin{bmatrix} b_x^2 & 0 \\ 0 & b_y^2 \end{bmatrix} \rightarrow \text{norm to 1}$$

μ doesn't matter (\rightarrow normalize all filters to 1!), but Σ does! \oplus or $\boxed{\cdot}$?

$b = 1$ pixel \rightarrow $\boxed{\cdot}$

Efficiency: $\boxed{\cdot} * \text{img} = (\oplus * \text{img}) * \boxed{\cdot}$ 8 vs. 6 multiplications!

dot product: $\langle a, b \rangle = a \cdot b = a_1 b_1 + a_2 b_2 + \dots + a_n b_n = \|a\| \|b\| \cos \varphi$

φ is the similarity, measure by norm to 1 a and b

 applying a filter can be seen as taking dot product \rightarrow find structures
this filter finds circles

$$\text{NCC: Normalized cross-correlation } \text{NCC}(I, T) = \frac{1}{n} \sum_{x,y} \frac{(I(x, y) - \bar{I})(T(x, y) - \bar{T})}{\delta_I \delta_T}$$

\bar{I} : mean of img, \bar{T} : mean of template

δ_I, δ_T : stdv of intensities \rightarrow scale to $\in [-1, 1]$ (remove mean: all at 0)

• no longer linear! $\delta^2 = \frac{1}{n} (\sum x^2) - \mu^2$

• intensity invariant (find face in highly illuminated pic is no problem)

size of filter matters \rightarrow change img (not filter) size

first filter coarse scales (gross) \rightarrow then refine, have different blur amounts

Gaussian pyramid: create each level from previous one: smooth & sample
gaussians are low pass filters

Summary: linear filtering: Form a new image whose pixels are a weighted sum of original pixel values

e.g. box filter / gaussian filter smoothing,
searching for a template

properties: output is a shift-invariant function of the input (same at each img location)

pyramid representations: important for describing and searching an img at all scales.

smooth to avoid aliasing! (\rightarrow vgl. aliasing)

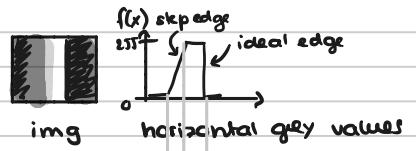
LEC 4 - Edges, corners & lines

L4 - 1

Edges contain most shape information!

- Derivatives & edges: detect changes (edges) with derivatives:

$$\frac{df}{dx} \approx [f(x+1) - f(x-1)]/2 \quad \text{img}$$



$$\text{gradient vector: } \left[\begin{array}{c} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{array} \right] \quad \approx f(x+1) - f(x-1)$$

$\begin{bmatrix} 1 & 0 & -1 \end{bmatrix}$ and $\begin{bmatrix} 1 \\ 0 \\ -1 \end{bmatrix}$ horizontal / vertical edged. filt.

combined with smoothing: Prewitt filters: $P_x = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$ $P_y = \begin{bmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{bmatrix}$

magnitude (stärke) of gradient: $\sqrt{P_x^2 + P_y^2}$ or $\max\{|P_x|, |P_y|\}$

- causes for edges: anisotropic (verde-Hintergrund), 3D shape , reflectance change (e.g. surface material change), illumination discontinuity (shadow)

- computational model: $\begin{bmatrix} -1 & 1 \end{bmatrix}$ is a convolution (approximation of $\frac{\partial f}{\partial x}$)

- noise: is detected by edge filters \rightarrow smooth before finding edges!
edge is location with high gradient, 2 derivatives: x and y direction
we can use derivatives of gaussian filters because:

$$D * (G * I) = (D * G) * I \quad D: \text{derivative}, G: \text{gauss smooth}, I: \text{img}$$

- Gradient magnitude: $I(x, y)$: img
 $I_x(x, y)$: der. in x direction = I_x
 $[I_x, I_y]$: vector gradient

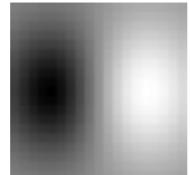
$$\sqrt{I_x^2 + I_y^2} = \text{gradient magnitude}$$

$$\begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

$$\begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix}$$

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 4 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$



derivative of Gauss in x direction

- more smoothing \rightarrow eliminates noise edges, makes edges smoother & thicker, removes fine detail

- canny edge detection:

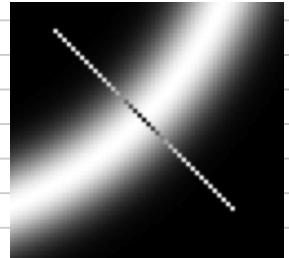
1.) Apply derivative of gaussian $(D * G) * I \rightarrow I_x, I_y$

2.) compute magnitude of gradients $\sqrt{I_x^2 + I_y^2}$

3.) Non-maximum suppression: thin multi-pixel-wide 'ridge' (= kante) down to single pixel

4.) link & threshold:

- low and high edge thresholds \rightarrow accept one high- + also accept low tr. ones that are connected to high threshold! \hookrightarrow Track edge by hysteresis



- hysteresis: "depending on its history", a lag or momentum factor

idea: 2 thresholds: k_{high} and k_{low}

use k_{high} to find strong edges to start edge chain

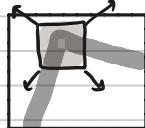
use k_{low} to find weak edges which continue edge chain

typical threshold ratio $k_{high}/k_{low} = 2$

\rightarrow closes gaps! — — \Rightarrow —

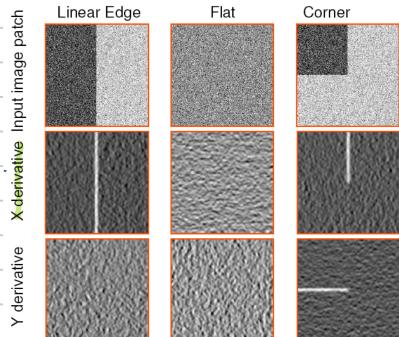
corners: edge detectors perform poorly at corners, but are important to detect

- at corners, gradient is ill defined, but around corner: 2 directions



- Harris corner: shuffle
 - flat region: no change
 - edge: change in 1 direction, no change in "edge" direction.
 - corner: change in all directions

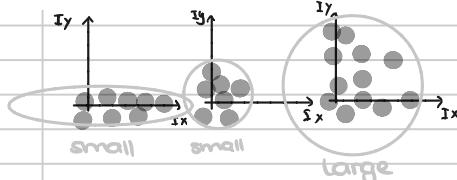
$$SSD(u, v) = \sum_{x,y} I(x+u, y+v) - I(x, y))^2 = [u \ v] \begin{bmatrix} \Sigma I_x^2 & \Sigma I_x I_y \\ \Sigma I_x I_y & \Sigma I_y^2 \end{bmatrix} [u \ v]$$



- corner detection must be rotation invariant! \rightarrow Eigenvalues

"principal component ellipse" which is only large at edges

- \downarrow data \rightarrow center data around 0
 \rightarrow compute covariance = inner product
 Eigenvalues tell us "change" in data, the "directions"



- Eigenvalue problem: λ : scalar = Eigenvalue

\vec{v} : Eigenvector, C : Matrix

$$C\vec{v} = \lambda \vec{v}$$

$$\Leftrightarrow \|C - \lambda I\| = 0 \quad I = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \dim(C) = 3,3$$

$$\left\| \begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix} - \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} \right\| = 0 = \left\| \begin{bmatrix} C_{11} - \lambda_1 & C_{12} \\ C_{21} & C_{22} - \lambda_2 \end{bmatrix} \right\| = 0$$

solve for λ_1, λ_2 and get \vec{v}_1 and \vec{v}_2

\Rightarrow Eigenvalues not affected by rotation! corner: λ_1, λ_2 are both large!

- computational considerations: $R = \det M - k(\text{trace } M)^2$ M : covariance mat
 det is connected to eigenvalues
- detect flat regions: below threshold: edge, within: flat, above: corner ?

Harris Corner Detection summary

convolution, point product
 is fast

- Compute x and y derivatives of image

convolve gradients

$$I_x = G_x^\sigma * I \quad I_y = G_y^\sigma * I$$

gaussian smoothed edges, 2. how much smoothing?

"compute covariance matrix"

- Compute product of derivatives at each pixel

point product

$$I_{x^2} = I_x I_x \quad I_{xy} = I_x I_y \quad I_{y^2} = I_y I_y$$

I_{xy} is an image

use I_x, I_y from step 1 and get

- Compute sums of products of derivatives at each pixel

conv

$$S_{x^2} = 1_{|I|} * I_{x^2} \quad S_{xy} = 1_{|I|} * I_{xy} \quad S_{y^2} = 1_{|I|} * I_{y^2}$$

$I_{x^2} * \dots = s_x^2$
 sum of values 1 of size t

- Define at each pixel the matrix

slow

$$H(x, y) = \begin{bmatrix} S_{x^2}(x, y) & S_{xy}(x, y) \\ S_{xy}(x, y) & S_{y^2}(x, y) \end{bmatrix}$$

- Compute the response of the detector at each location

$$R = \text{Det}(H) - k(\text{Trace}(H))^2$$

- Threshold on the value of R. Compute non-maximum suppression.

depending on whether you want to detect flats, edges or corners